

Reinforcement Learning for Large Language Models

Making the Hard Math Easy

TABLE OF CONTENTS

PART 1 FOUNDATIONS

1. Essential Math Toolkit
2. Why LLMs Need RL: The Alignment Gap
3. RL Fundamentals: The Complete Picture
4. Setting Up Your Free Training Environment

PART 2 CORE METHODS

5. Supervised Fine-Tuning & The Cold Start
6. RLHF: The Three-Step Dance
7. Direct Preference Optimization
8. Online DPO & Iterative Alignment
9. Reward Modeling & The Critic
10. Modern RL Algorithms: GRPO, RLOO, KTO & More

PART 3 ADVANCED TECHNIQUES

11. Verifiers & Outcome Rewards: Beyond PRMs
12. Reasoning with GRPO: The DeepSeek Recipe
13. Test-Time Compute: Scaling at Inference
14. Self-Play & Constitutional AI
15. Multi-Objective RL & Agentic Systems
16. Domain-Specific RL: Code, Math, Tools, Dialogue

PART 4 PRODUCTION & RECIPES

17. Three Recipes: Chatbot, Reasoner, Agent
18. Debugging, Scaling & Deployment
19. Your RL Journey: From Here to Production